
Assignment 1 - Exercise 3

Table of Contents

(2)	1
(4 & 5)	3
"Gradient Descent"	13
"Conjugate gradient"	14

Name: Stefano Gonçalves Simao

Date: 26/3/2021

(2)

```
close all;
clc;

disp('(2)');
figure(1);
h1 = subplot (2,2,1);
mu = 1;
[x,y] = meshgrid(-10:0.5:10, -10:0.5:10);
f = x.^2 + mu * y.^2;
surf(x,y,f)
title('Surface with \mu = 1');

h2 = subplot (2,2,2);
mu = 10;
[x,y] = meshgrid(-10:0.5:10, -10:0.5:10);
f = x.^2 + mu * y.^2;
surf(x,y,f)
title('Surface with \mu = 10');

h3 = subplot (2,2,3);
mu = 1;
[x,y] = meshgrid(-10:0.5:10, -10:0.5:10);
f = x.^2 + mu * y.^2;
contour(x,y,f)
title('Contour with \mu = 1');

h4 = subplot (2,2,4);
mu = 10;
[x,y] = meshgrid(-10:0.5:10, -10:0.5:10);
f = x.^2 + mu * y.^2;
contour(x,y,f)

title('Contour with \mu = 10');

disp('-----')
```

```

disp('The starting points in the function with  $\mu = 10$  are the ones
    that could be problematic.')
disp("It's much difficult to find a good starting point and it should
    be where the contour lines")
disp("meet the axes as the gradient of a function is always
    perpendicular to the")
disp("contour lines (a good place would be with any y but around x =
    0). Here the convergence will")
disp("be slower as it will zigzag much more. This is beacause the
    contour lines are ellipses whose")
disp('axes lie along the orthogonal lines of A, and with this
    difference in magnitude, as said,')
disp('the steepest descent will oscillate.')
disp("In the other case it's much easier for the algorithm to converge
    as every starting point has")
disp("fast convergence. It should be able to converge in 1
    iteration.")
disp('-----')
disp(' ')

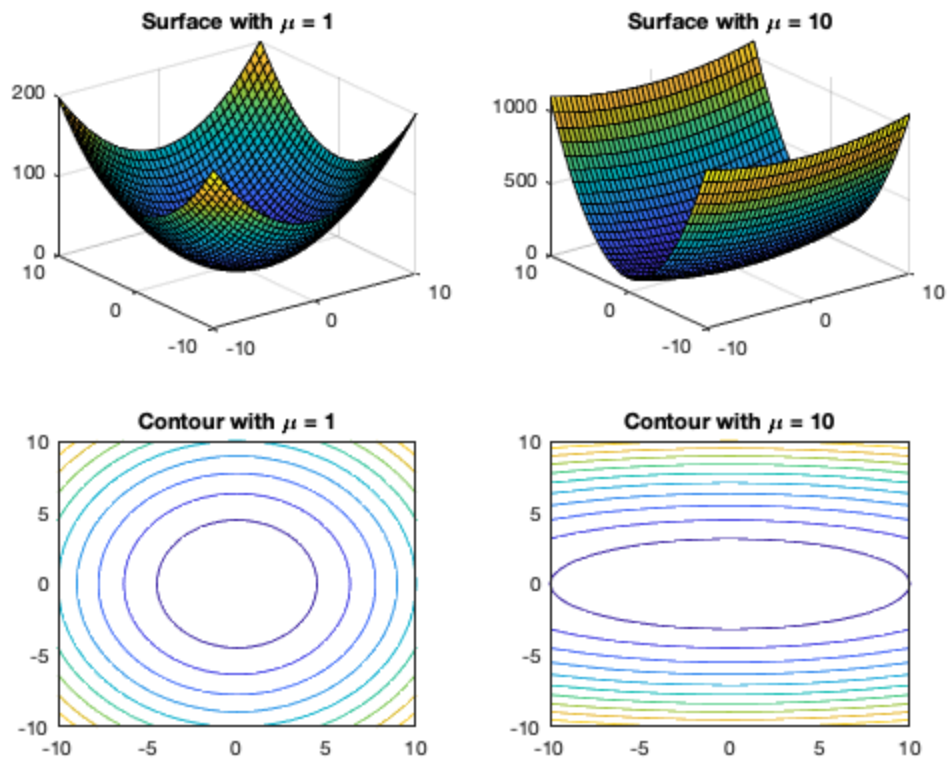
```

(2)

```

-----
The starting points in the function with  $\mu = 10$  are the ones that
    could be problematic.
It's much difficult to find a good starting point and it should be
    where the contour lines
meet the axes as the gradient of a function is always perpendicular to
    the
contour lines (a good place would be with any y but around x = 0).
    Here the convergence will
be slower as it will zigzag much more. This is beacause the contour
    lines are ellipses whose
axes lie along the orthogonal lines of A, and with this difference in
    magnitude, as said,
the steepest descent will oscillate.
In the other case it's much easier for the algorithm to converge as
    every starting point has
fast convergence. It should be able to converge in 1 iteration.
-----

```



(4 & 5)

```
disp('(4 & 5)')
disp('-----')
disp('Not knowing exactly what is required in the assignment, I decided
to implement both Gradient descent')
disp('and Conjugate Gradient methods. Setting beta = 0 in the CD
algorithm would make x computed by the')
disp('gradient descent method. One thing to note here is that in the
plots that are functions of')
disp('the iterations, the iteration starts at number 2, the iteration
No.1 is the status before ')
disp('the start of the loop.')
disp('-----')
%Gradient descent first
figure('Name', 'Gradient descent');
v = tiledlayout(2,3);
title(v, 'Gradient descent')
[x1, r1, F1]=GD(1, [10;0]);
[x2, r2, F2]=GD(1, [0;10]);
[x3, r3, F3]=GD(1, [10;10]);

[x1, r4, F4]=GD(10, [10;0]);
[x2, r5, F5]=GD(10, [0;10]);
[x3, r6, F6]=GD(10, [10;10]);
```

```
%Gradient norm
figure('Name', 'Gradient descent - log10 norm of the Gradient');
z = tiledlayout(2,3);
nexttile;
semilogy(r1, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 1 and x0 = (10,0)');
title(s)

nexttile;
semilogy(r2, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 1 and x0 = (0,10)');
title(s)

nexttile;
semilogy(r3, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 1 and x0 = (10,10)');
title(s)

nexttile;
semilogy(r4, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 10 and x0 = (10,0)');
title(s)

nexttile;
semilogy(r5, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 10 and x0 = (0,10)');
title(s)

nexttile;
semilogy(r6, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 10 and x0 = (10,10)');
title(s)
title(z, 'Gradient descent - Energy function')

%Energy function value
figure('Name', 'Gradient descent - Energy function');
j = tiledlayout(2,3);
nexttile;
plot(F1, '-go');
xlabel('Iteration');
ylabel('Function value');
```

```

s = sprintf('μ = 1 and x0 = (10,0)');
title(s)

nexttile;
plot(F2, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 1 and x0 = (0,10)');
title(s)

nexttile;
plot(F3, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 1 and x0 = (10,10)');
title(s)

nexttile;
plot(F4, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 10 and x0 = (10,0)');
title(s)

nexttile;
plot(F5, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 10 and x0 = (0,10)');
title(s)

nexttile;
plot(F6, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 10 and x0 = (10,10)');
title(s)
title(j,'Gradient descent - Energy function')

disp('-----')
%Conjugate gradient second
figure('Name', 'Conjugate gradient');
t = tiledlayout(2,3);
title(t,'Conjugate gradient')
[x1, r1, F1]=CG(1, [10;0]);
[x2, r2, F2]=CG(1, [0;10]);
[x3, r3, F3]=CG(1, [10;10]);

[x1, r4, F4]=CG(10, [10;0]);
[x2, r5, F5]=CG(10, [0;10]);
[x3, r6, F6]=CG(10, [10;10]);

%Gradient norm
figure('Name', 'Conjugate gradient - log10 norm of the Gradient');
```

```
g = tiledlayout(2,3);
nexttile;
semilogy(r1, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 1 and x0 = (10,0)');
title(s)

nexttile;
semilogy(r2, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 1 and x0 = (0,10)');
title(s)

nexttile;
semilogy(r3, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 1 and x0 = (10,10)');
title(s)

nexttile;
semilogy(r4, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 10 and x0 = (10,0)');
title(s)

nexttile;
semilogy(r5, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 10 and x0 = (0,10)');
title(s)

nexttile;
semilogy(r6, '-r+');
xlabel('Iteration');
ylabel('Gradient norm');
s = sprintf('μ = 10 and x0 = (10,10)');
title(s)
title(g, 'Conjugate gradient - log10 norm of the Gradient')

%Energy function value
figure('Name', 'Conjugate gradient - Energy function');
u = tiledlayout(2,3);
nexttile;
plot(F1, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 1 and x0 = (10,0)');
title(s)
```

```

nexttile;
plot(F2, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 1 and x0 = (0,10)');
title(s)

nexttile;
plot(F3, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 1 and x0 = (10,10)');
title(s)

nexttile;
plot(F4, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 10 and x0 = (10,0)');
title(s)

nexttile;
plot(F5, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 10 and x0 = (0,10)');
title(s)

nexttile;
plot(F6, '-go');
xlabel('Iteration');
ylabel('Function value');
s = sprintf('μ = 10 and x0 = (10,10)');
title(s)
title(u,'Conjugate gradient - Energy function')

disp('-----')
disp("We can see that with μ = 1 there are no problems, each algorithm
    finds the solution in");
disp("one iteration as the negative gradient points to the solution.
    With μ = 10 it's more interesting");
disp("as for (x0,y0) = (10,0) and (0,10) the starting position is a
    good choice for both");
disp("algorithms. (10,10) is trickier here as in GD it zigzags more as
    it is not a good ");
disp("starting point. With CG we are sure to come to the solution in n
    iterations, in this");
disp("case we have n = 2. So it is anyway more effective even with a
    starting point that is not good.");
disp("In the case of μ = 10 and starting point (10,10), both
    algorithms don't provide an exact");
disp("solution, but one that is very close to (0,0). It is inside the
    tol given.");

```

```

disp("Another thing to note is that in the log10 of the norm of the
    gradient ");
disp("the zero is not represented as log(0) = Inf for Matlab and it
    doesn't plot it.");
disp('-----')

(4 & 5)
-----
Not knowing exactly what is required in the assignment, I decided to
    implement both Gradient descent
    and Conjugate Gradient methods. Setting beta = 0 in the CD algorithm
    would make x computed by the
    gradient descent method. One thing to note here is that in the plots
    that are functions of
    the iterations, the iteration starts at number 2, the iteration No.1
    is the status before
    the start of the loop.
-----
GD
     $\mu = 1$ : and starting point (10,0)
x =
    0
    0

-----
GD
     $\mu = 1$ : and starting point (0,10)
x =
    0
    0

-----
GD
     $\mu = 1$ : and starting point (10,10)
x =
    0
    0

-----
GD
     $\mu = 10$ : and starting point (10,0)
x =
    0
    0

-----
GD
     $\mu = 10$ : and starting point (0,10)
x =
    0
    0

-----
GD

```



```

 $\mu = 10$ : and starting point (10,10)
x =
    1.0e-08 *

    0.7710
   -0.0077

```

CG

```

 $\mu = 1$ : and starting point (10,0)
x =
     0
     0

```

CG

```

 $\mu = 1$ : and starting point (0,10)
x =
     0
     0

```

CG

```

 $\mu = 1$ : and starting point (10,10)
x =
     0
     0

```

CG

```

 $\mu = 10$ : and starting point (10,0)
x =
     0
     0

```

CG

```

 $\mu = 10$ : and starting point (0,10)
x =
     0
     0

```

CG

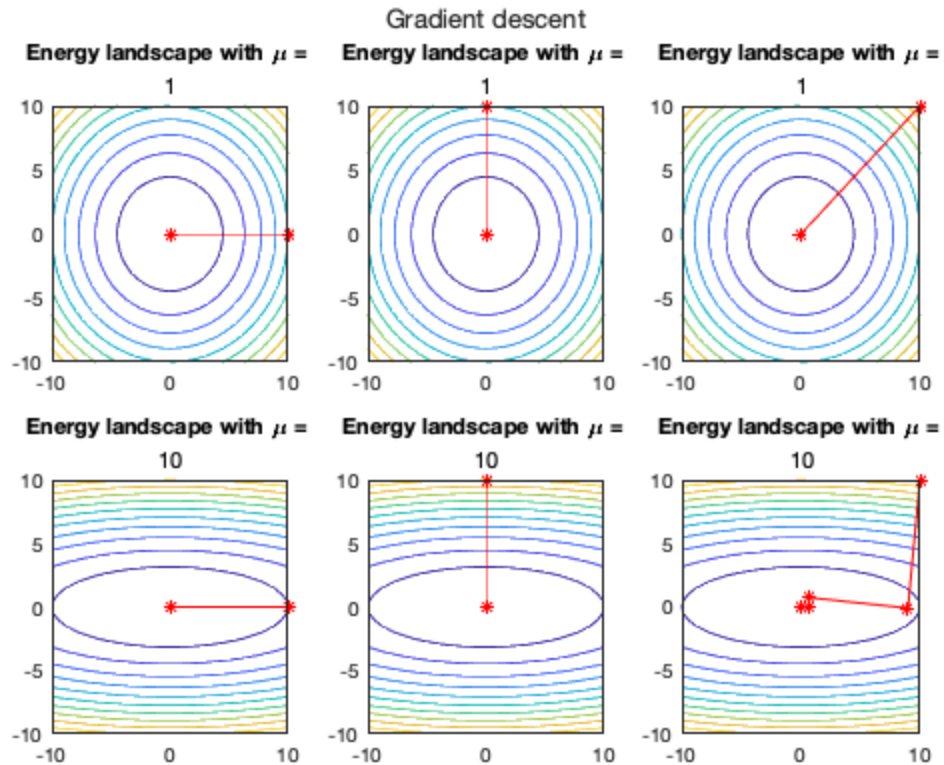
```

 $\mu = 10$ : and starting point (10,10)
x =
    1.0e-14 *

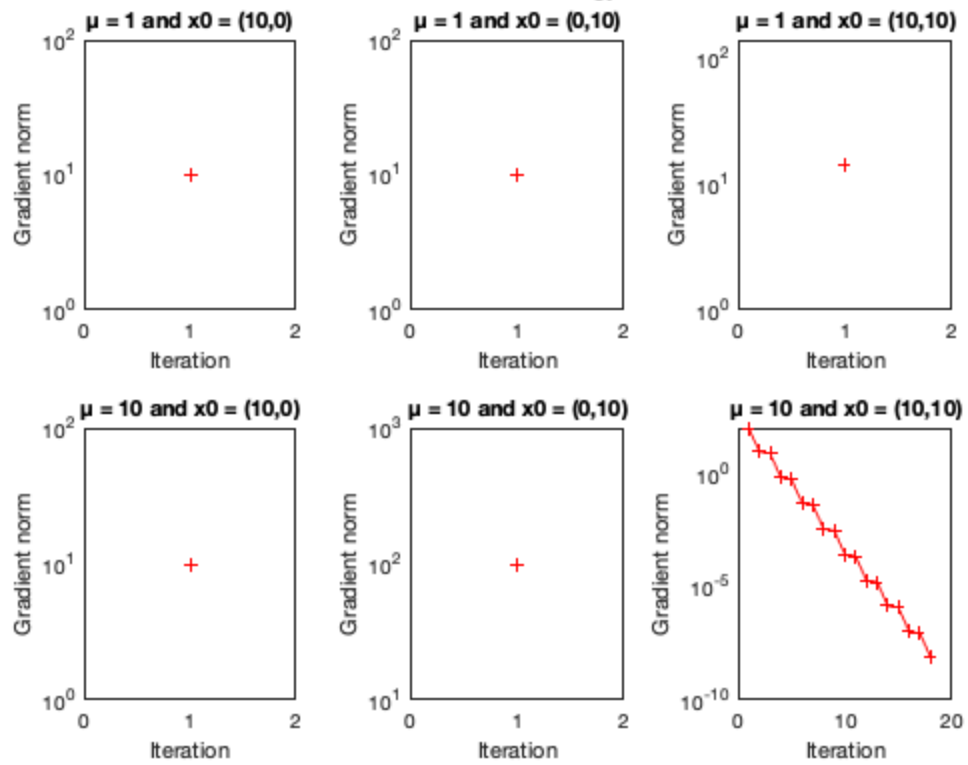
   -0.1776
    0.2012

```

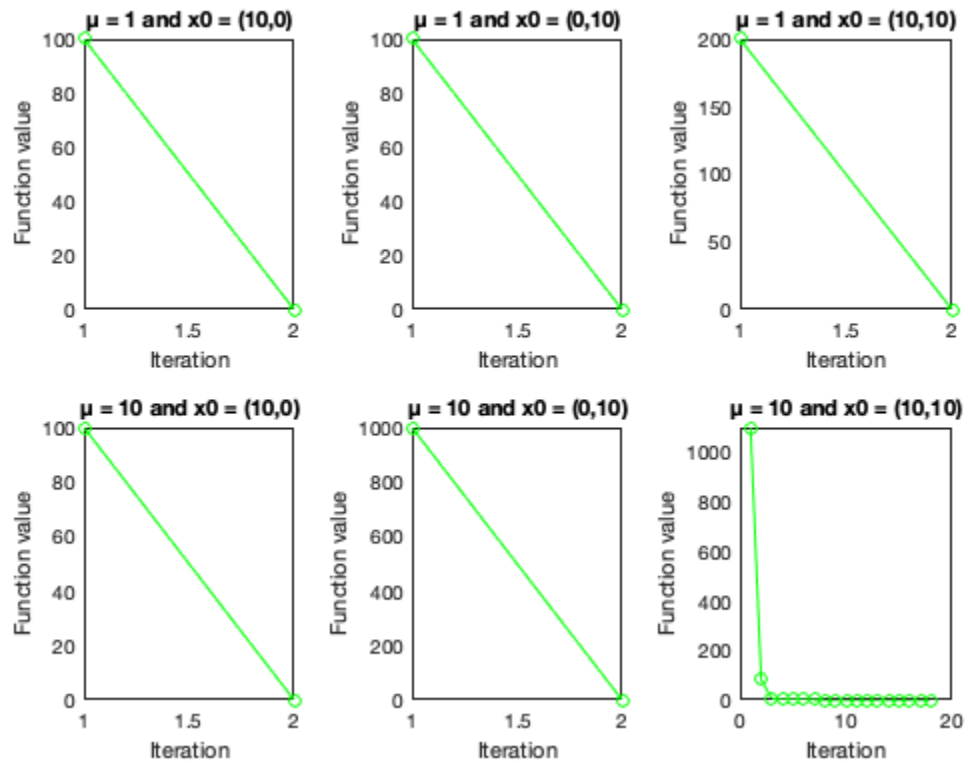
We can see that with $\mu = 1$ there are no problems, each algorithm finds the solution in one iteration as the negative gradient points to the solution. With $\mu = 10$ it's more interesting as for $(x_0, y_0) = (10, 0)$ and $(0, 10)$ the starting position is a good choice for both algorithms. $(10, 10)$ is trickier here as in GD it zigzags more as it is not a good starting point. With CG we are sure to come to the solution in n iterations, in this case we have $n = 2$. So it is anyway more effective even with a starting point that is not good. In the case of $\mu = 10$ and starting point $(10, 10)$, both algorithms don't provide an exact solution, but one that is very close to $(0, 0)$. It is inside the tol given. Another thing to note is that in the log10 of the norm of the gradient the zero is not represented as $\log(0) = \text{Inf}$ for Matlab and it doesn't plot it.

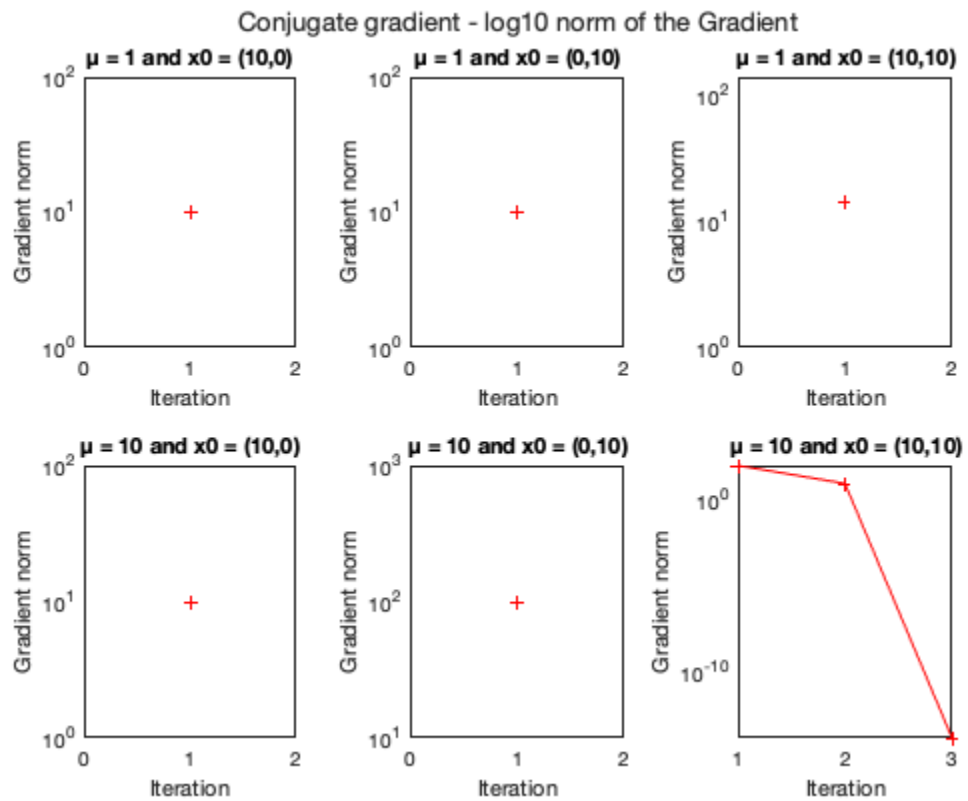
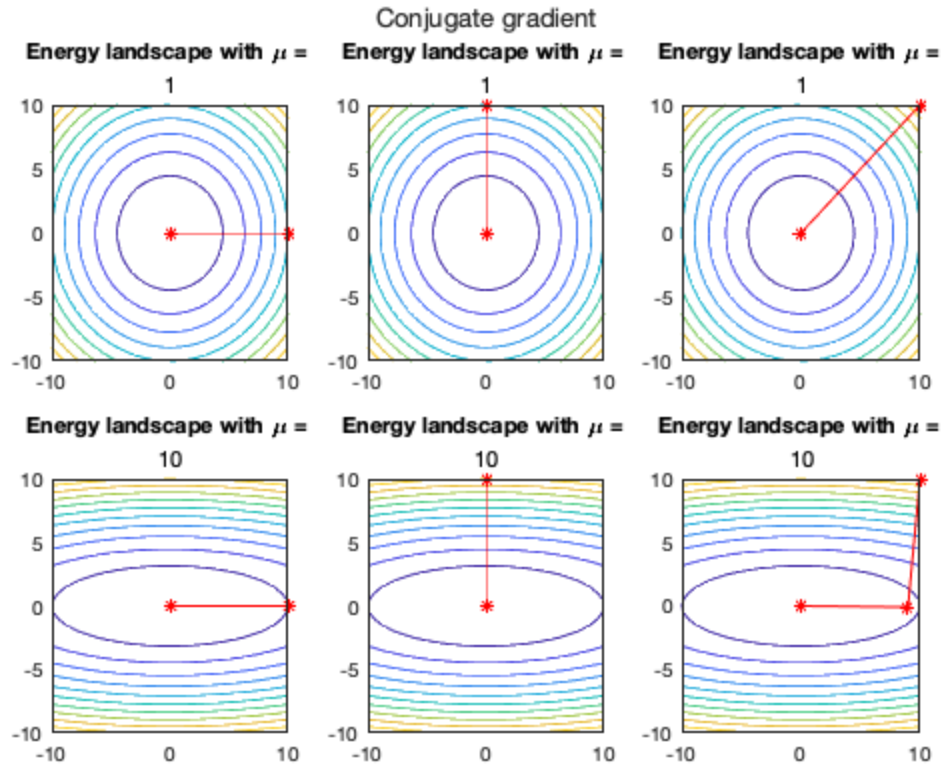


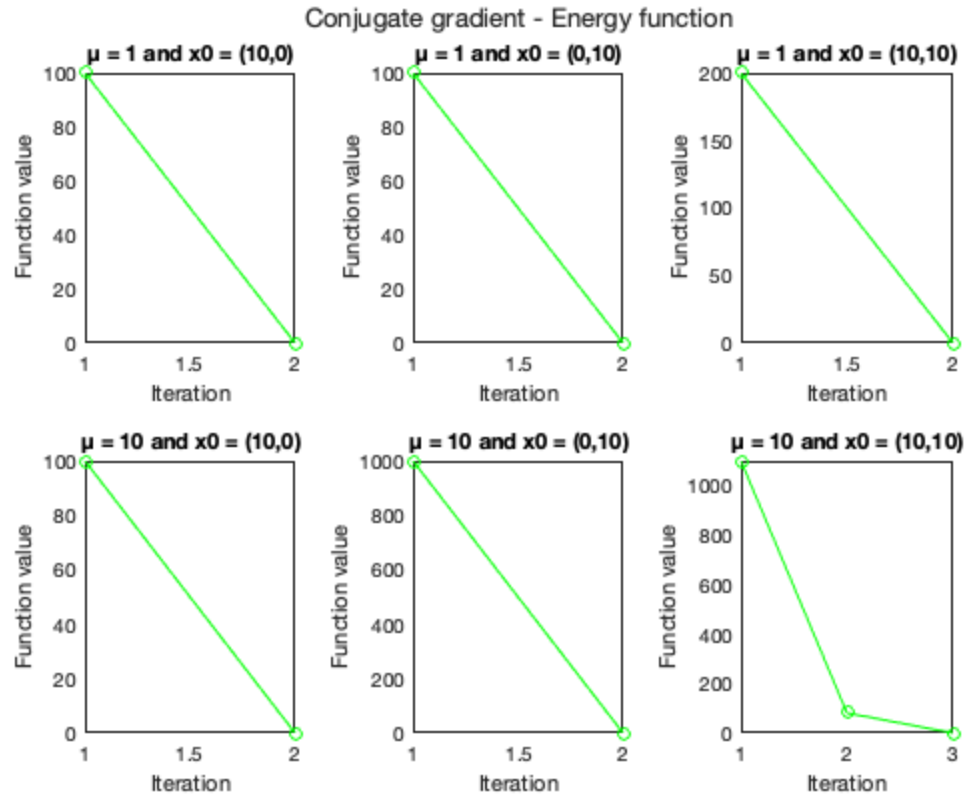
Gradient descent - Energy function



Gradient descent - Energy function







"Gradient Descent"

```
function [x, rvec, F]=GD(mu,x0)
%Gradient descent method

%Initialization
A = [1 0; 0 mu];
b = [0; 0];
x = x0;
rvec = [];
Xarr = [];
Yarr = [];
F = [];
max_itr = 100;
tol = 10^(-8);

nexttile;
formatSpec = 'GD\n \mu = %d: and starting point (%d,%d)\n';
fprintf(formatSpec, mu, x0(1,1), x0(2,1));
[X,Y] = meshgrid(-10:0.5:10, -10:0.5:10);
f = X.^2 + mu * Y.^2;
contour(X,Y,f)
title('Energy landscape with \mu = ', num2str(mu));
hold on;
```

```

r = b - A * x0;
d = r;
p_old = dot(r,r);
i = 1;
%The first iteration is actually the status before it starts.
    Iteration 1
%starts at position 2
rvec(i) = norm(r);
F(i) = x0(1,1)^2 + mu * x0(2,1)^2;

Xarr(1) = x0(1,1);
Yarr(1) = x0(2,1);

while (i < max_itr && norm(r) >= tol)
    s = A * d;
    alpha = p_old / dot(d, s);
    x = x + alpha * d;
    r = r - alpha * s;
    p_new = dot(r,r);
    beta = 0;
    d = r + beta * d;
    p_old = p_new;
    i = i + 1;
    Xarr(end + 1) = x(1,1);
    Yarr(end + 1) = x(2,1);
    rvec(i) = norm(r);
    F(i) = x(1,1)^2 + mu * x(2,1)^2;
end
plot (Xarr,Yarr, '-r*');
disp('x =');
disp(x);
disp('-----');
end

```

"Conjugate gradient"

```

function [x, rvec, F]=CG(mu,x0)
%Conjugate Gradient method

%Initialization
A = [1 0; 0 mu];
b = [0; 0];
x = x0;
rvec = [];
Xarr = [];
Yarr = [];
F = [];
max_itr = 100;
tol = 10^(-8);

```

```

nexttile;
formatSpec = 'CG\n  $\mu$  = %d: and starting point (%d,%d)\n';
fprintf(formatSpec, mu, x0(1,1), x0(2,1));
[X,Y] = meshgrid(-10:0.5:10, -10:0.5:10);
f = X.^2 + mu * Y.^2;
contour(X,Y,f)
title('Energy landscape with  $\mu$  = ', num2str(mu));
hold on;

r = b - A * x0;
d = r;
p_old = dot(r,r);
i = 1;
%The first iteration is actually the status before it starts.
    Iteration 1
%starts at position 2
rvec(i) = norm(r);
F(i) = x0(1,1)^2 + mu * x0(2,1)^2;

Xarr(1) = x0(1,1);
Yarr(1) = x0(2,1);

while (i < max_itr && norm(r) > tol)
    s = A * d;
    alpha = p_old / dot(d, s);
    x = x + alpha * d;
    r = r - alpha * s;
    p_new = dot(r,r);
    beta = p_new/p_old;
    d = r + beta * d;
    p_old = p_new;
    i = i + 1;
    Xarr(end + 1) = x(1,1);
    Yarr(end + 1) = x(2,1);
    rvec(i) = norm(r);

    F(i) = x(1,1)^2 + mu * x(2,1)^2;
end
plot (Xarr,Yarr, '-r*');
disp('x = ');
disp(x);
disp('-----');
end

```

Published with MATLAB® R2020b